

# 你必须知道的技巧

王树森


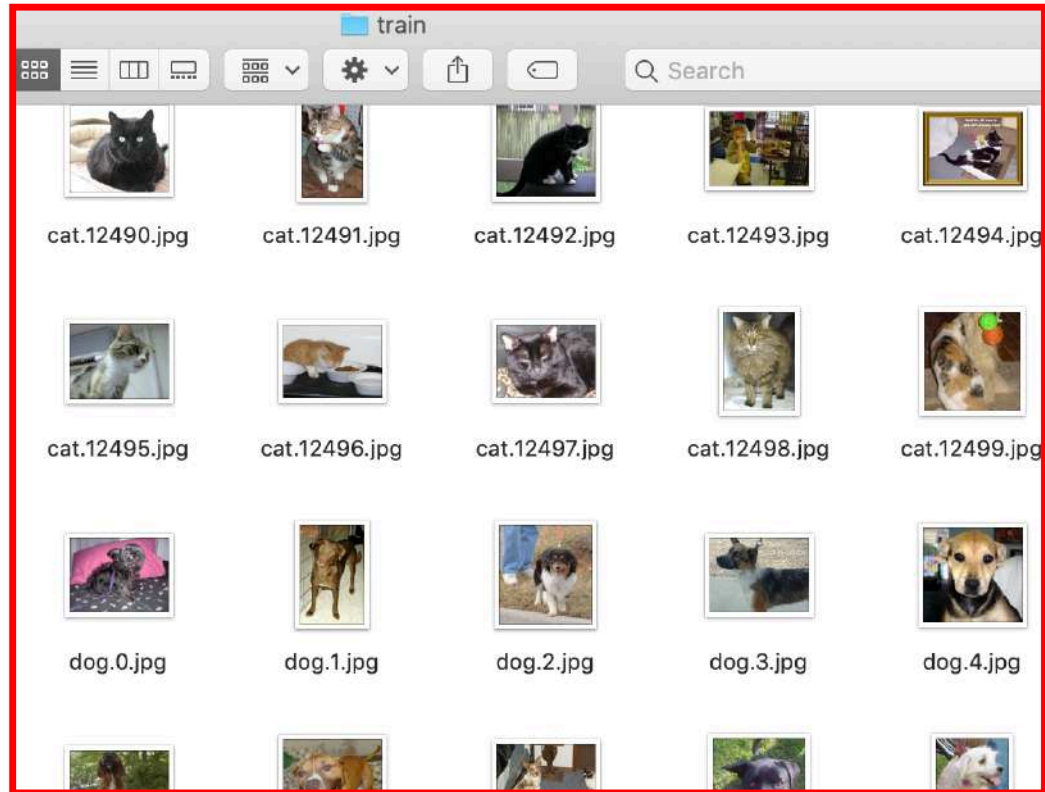
# **The Dogs vs. Cats Dataset**

# The Dogs vs. Cats Dataset




# The Dogs vs. Cats Dataset

Download: <https://www.kaggle.com/c/dogs-vs-cats/data>




# The Dogs vs. Cats Dataset

Download: <https://www.kaggle.com/c/dogs-vs-cats/data>



cat.0.jpg




cat.1.jpg



cat.2.jpg



cat.3.jpg



cat.4.jpg



cat.5.jpg



cat.6.jpg



cat.7.jpg




cat.8.jpg




cat.9.jpg



cat.10.jpg



cat.11.jpg



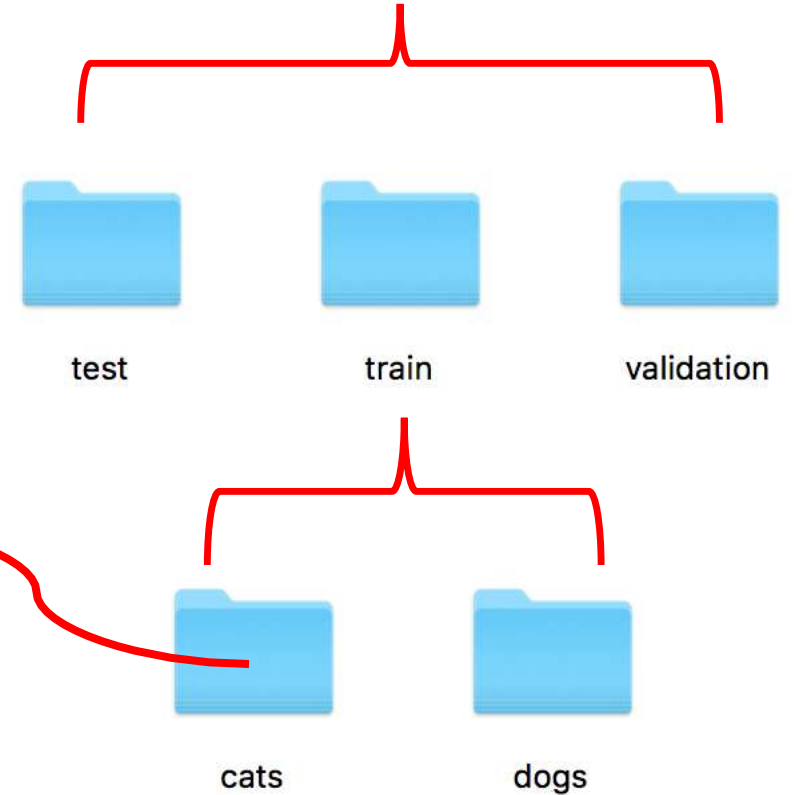
cat.12.jpg



cat.13.jpg



cat.14.jpg



# The Dogs vs. Cats Dataset

- 数据集有 25000 个样本.
- 使用一个子集:
  - 2000 images for 训练,
  - 1000 images for 验证,
  - 1000 images for 测试.

用 keras 实现一个CNN

# 1. Load and Process the Dataset

- 当前，文件以JPEG格式呈现
- 数据处理：
  1. 读取数据文件
  2. 把JPEG文件解码成三阶张量
  3. 把文件调整为相同大小  $150 \times 150 \times 3$
  4. 将像素值（在 0 到 255 之间）重新缩放到  $[0, 1]$  区间



# 1. Load and Process the Dataset

```
from keras.preprocessing.image import ImageDataGenerator
```

```
# All images will be rescaled by 1./255
```

```
train_datagen = ImageDataGenerator(rescale=1./255)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

# 1. Load and Process the Dataset

```
from keras.preprocessing.image import ImageDataGenerator

# All images will be rescaled by 1./255
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)



train_generator = train_datagen.flow_from_directory(
    # This is the target directory
    train_dir,
    # All images will be resized to 150x150
    target_size=(150, 150),
    batch_size=20,
    # Since we use binary_crossentropy loss, we need binary labels
    class_mode='binary')
```

# 1. Load and Process the Dataset


```
for data_batch, labels_batch in train_generator:  
    print('data batch shape:', data_batch.shape)  
    print('labels batch shape:', labels_batch.shape)  
    break
```

```
data batch shape: (20, 150, 150, 3)  
labels batch shape: (20,)
```

# 1. Load and Process the Dataset



# 1. Load and Process the Dataset



## 2. Build the CNN

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

## 2. Build the CNN

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

### 3. Train the CNN

指定：优化方法、学习率（LR）、损失函数和评估指标。

```
from keras import optimizers


model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```



### 3. Train the CNN

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50)
```

### 3. Train the CNN

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,   
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50)
```

- Totally  $n = 2000$  training samples.
- Batch size is  $b = 20$ .
- Thus  $\frac{n}{b} = 100$  batches per epoch.


```
Epoch 1/30  
100/100 [=====] - 9s - loss: 0.6898 - acc: 0.5285 - val_loss: 0.6724 - val_acc: 0.5950  
Epoch 2/30  
100/100 [=====] - 8s - loss: 0.6543 - acc: 0.6340 - val_loss: 0.6565 - val_acc: 0.5950  
Epoch 3/30  
100/100 [=====] - 8s - loss: 0.6143 - acc: 0.6690 - val_loss: 0.6116 - val_acc: 0.6650  
Epoch 4/30  
100/100 [=====] - 8s - loss: 0.5626 - acc: 0.7125 - val_loss: 0.5774 - val_acc: 0.6970
```



```
Epoch 29/30  
100/100 [=====] - 8s - loss: 0.0375 - acc: 0.9915 - val_loss: 0.9987 - val_acc: 0.7100  
Epoch 30/30  
100/100 [=====] - 8s - loss: 0.0387 - acc: 0.9895 - val_loss: 1.0139 - val_acc: 0.7240
```

## 4. Examine the Results


Plot the **accuracy** against **epochs** (1 epoch = 1 pass over the data).



- 训练准确率: 99.0%
- 测试准确率: 72.4%
- 好像过拟合了

## 4. Examine the Results

Plot the **loss** against **epochs** (1 epoch = 1 pass over the data).



- 训练损失函数持续下降
- 验证损失函数下降后上升

# Why Overfitting?

=====

Total params: 3,453,121

Trainable params: 3,453,121

Non-trainable params: 0

---


超过  $3M$  个参数; 但是只有  $2K$  个训练样本.  
过拟合并不意外

# **Trick 1: Dropout**


# Dropout: 基本概念

- Train

- 在训练的每次迭代（1 次前向 + 1 次后向）中，随机屏蔽 50%（或任意百分比）的神经元。



(a) Standard Neural Net



(b) After applying dropout.


# Dropout: 基本概念

- Train
  - 在训练的每次迭代（1 次前向 + 1 次后向）中，随机屏蔽 50%（或任意百分比）的神经元。
- Prediction
  - 不使用 dropout
  - 使用所有的参数




# Dropout: 实现

在这层执行dropout




- **For a batch of training samples**

- 随机选择50%的神经元
- 把这些神经元设置为0
- 将未被选中的神经元乘以  $\frac{1}{0.5} = 2$ .




# Dropout: 实现


在这层执行dropout



- **For a batch of training samples...**
- 对于另一个批次，进行一次独立的随机抽样（即随机屏蔽神经元）




# Dropout: 实现



- **Input:** vector  $\mathbf{x}^{(0)} \in \mathbb{R}^5$ .
- $\mathbf{z}^{(1)} = \mathbf{W}^{(0)} \mathbf{x}^{(0)} \in \mathbb{R}^4$ .
- $\mathbf{x}^{(1)} = \max\{\mathbf{0}, \mathbf{z}^{(1)}\} \in \mathbb{R}^4$ .
- $\mathbf{z}^{(2)} = \mathbf{W}^{(1)} \mathbf{x}^{(1)} \in \mathbb{R}^3$ .
- **Output:**  $\text{SoftMax}(\mathbf{z}^{(2)}) \in \mathbb{R}^3$ .

# Dropout: 实现

这层执行正则化  
(防止过拟合的技术就叫正则化)



- **Input:** vector  $\mathbf{x}^{(0)} \in \mathbb{R}^5$ .
- $\mathbf{z}^{(1)} = \mathbf{W}^{(0)} \mathbf{x}^{(0)} \in \mathbb{R}^4$ .
- $\mathbf{x}^{(1)} = \max\{\mathbf{0}, \mathbf{z}^{(1)}\} \in \mathbb{R}^4$ .
- Add a dropout layer
- $\mathbf{z}^{(2)} = \mathbf{W}^{(1)} \tilde{\mathbf{x}}^{(1)} \in \mathbb{R}^3$ .
- **Output:**  $\text{SoftMax}(\mathbf{z}^{(2)}) \in \mathbb{R}^3$ .

- $\mathbf{m} \in \mathbb{R}^4$  是一个随机向量 (Each entry is 0 or 2, w.p. 50%.)

- Apply  $\mathbf{m}$  to  $\mathbf{x}^{(1)}$ :  
 $\tilde{\mathbf{x}}^{(1)} = \mathbf{m} \circ \mathbf{x}^{(1)}$ .

“ $\circ$ ” is 逐元素乘法

# Keras's Dropout Layer

- 仅在第一个全连接层之前使用 **Dropout**，以正则化第一个全连接层。
- 因为第一个全连接层有太多的可训练参数。


```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```



# Why Does Dropout Work?

- 在训练中，Dropout 强制网络根据部分特征进行决策。



# Why Does Dropout Work?

- 在训练中，Dropout 强制网络根据部分特征进行决策。
- Dropout 是一种正则化机制 [1].
  - 缓解过拟合。
  - 类似于 L1 和 L2 范数正则化。
  - 但 Dropout 在经验上效果更好。

## Reference:

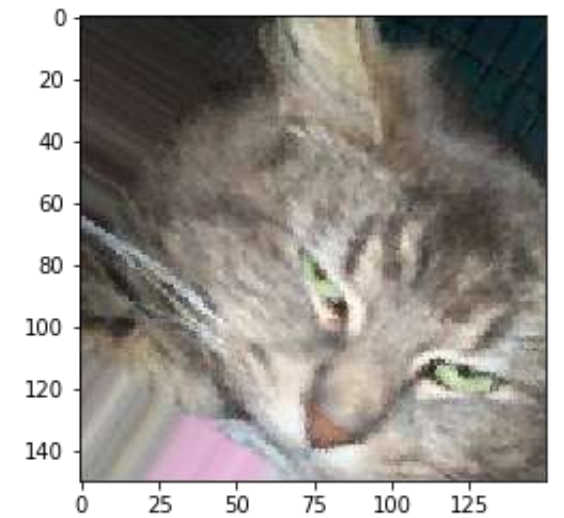



1. Wager, Wang, & Liang. Dropout Training as Adaptive Regularization. In *NIPS*, 2013.

## 技巧 2：数据增强



# 数据增强

- 数据增强：从现有训练数据生成更多训练样本。
- 例如，翻转、旋转、裁剪、平移、添加随机噪声。




# 数据增强：样本

水平翻转



# 数据增强：样本

随机裁剪和缩放



# 数据增强：样本

颜色抖动（随机调整对比度和亮度）



# Setup Data Augmentation Using Keras

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,)
```

```
train_generator = train_datagen.flow_from_directory(  
    # This is the target directory  
    train_dir,  
    # All images will be resized to 150x150  
    target_size=(150, 150),  
    batch_size=32,  
    # Since we use binary_crossentropy loss, we need binary labels  
    class_mode='binary')
```

# Train the CNN

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=100,  
    validation_data=validation_generator,  
    validation_steps=50)
```

Epoch 1/100

100/100 [=====] - 24s - loss: 0.6857 - acc: 0.5447 - val\_loss: 0.6620 - val\_acc: 0.5888

Epoch 2/100

100/100 [=====] - 23s - loss: 0.6710 - acc: 0.5675 - val\_loss: 0.6606 - val\_acc: 0.5825

Epoch 3/100

100/100 [=====] - 22s - loss: 0.6609 - acc: 0.5913 - val\_loss: 0.6663 - val\_acc: 0.5711

●  
●  
●

Epoch 99/100

100/100 [=====] - 22s - loss: 0.3255 - acc: 0.8581 - val\_loss: 0.3518 - val\_acc: 0.8460


Epoch 100/100

100/100 [=====] - 22s - loss: 0.3280 - acc: 0.8491 - val\_loss: 0.3776 - val\_acc: 0.8439




# Examine the Results

***accuracy*** against ***epochs***



***loss*** against ***epochs***



# 要点总结

- 要训练用于图像的卷积神经网络（ConvNet），始终使用数据增强.
- 它可以免费为你提供更多数据!
- 如果某一层有太多参数，在其之前添加一个 Gropout 层
- 正则化可以防止过拟合
- 训练速度会稍微变慢



## 技巧3:预训练

# Train a Deep Neural Network?

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```

- 我们已经训练了一个具有 4 个卷积层和 2 个全连接层的神经网络
- 相对较浅


# Train a Deep Neural Network?

我们可以训练一个深层神经网络吗？

- 难度很大
  - 训练的参数数量非常大
  - 深层网络表达能力非常强
  - 我们只有 2555 个训练样本
- 单纯地训练一个深层网络肯定会导致过拟合

解决方案：预训练


VGG16 网络



# 预训练


1. 在大规模数据集上预训练一个深层网络，例如，ImageNet（1400 万张带标签的图像）。

**VGG16 网络**



# 预训练

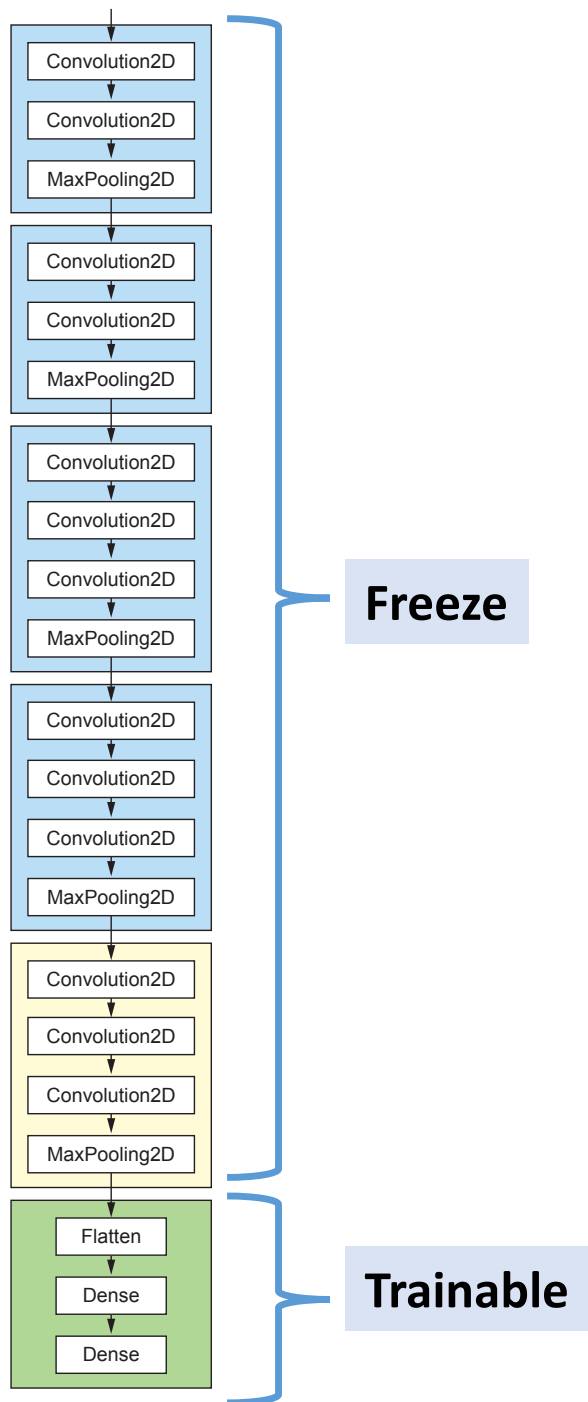
1. 在大规模数据集上预训练一个深层网络，例如，ImageNet（1155 万张带标签的图像）。
2. 为什么只移除最顶层？
  - 不同的输出形状和激活函数
  - 新的分类器专用于某种场景



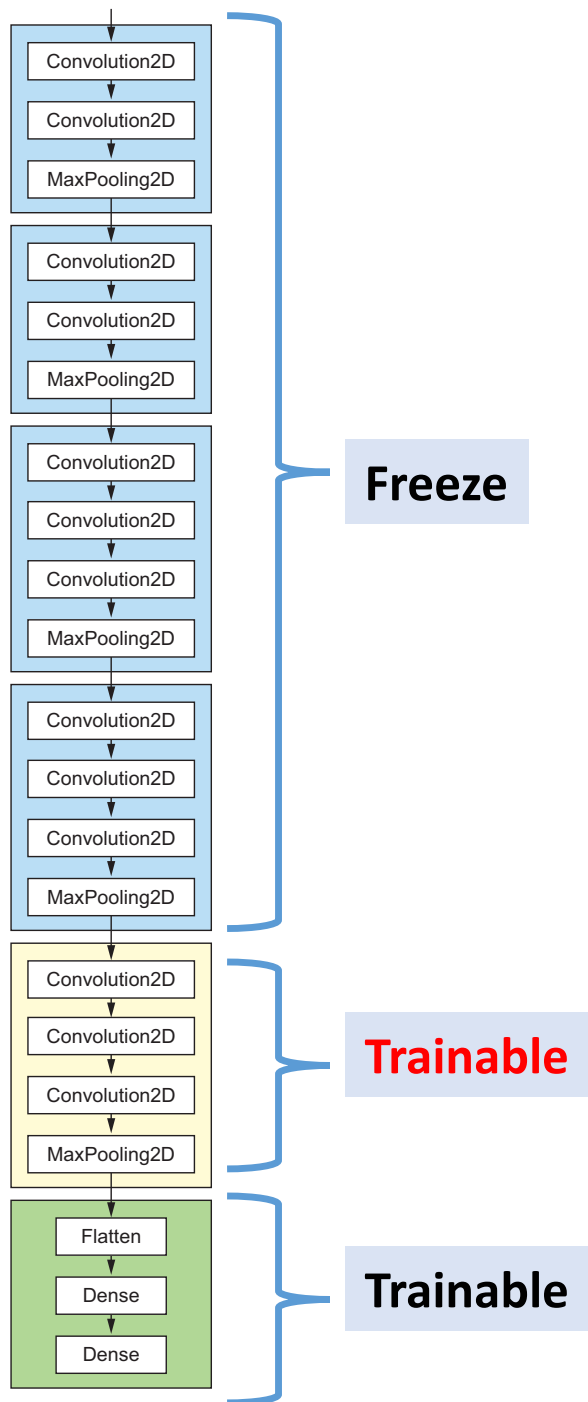
Remove the top layers

# 预训练

1. 在大规模数据集上预训练一个深层网络，例如，ImageNet（1400 万张带标签的图像）。
2. 移除最顶层
3. 搭建新的最顶层 (随机初始化).
4. 冻结其他层; 只训练最顶层

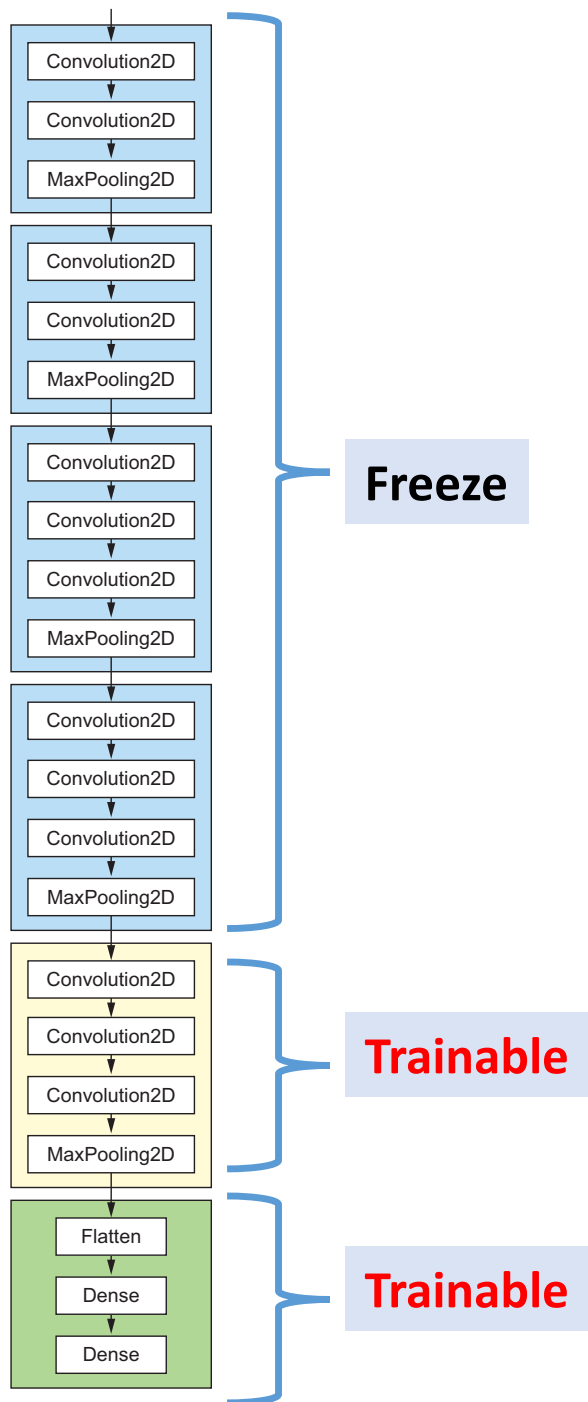


# 预训练



1. 在大规模数据集上预训练一个深层网络，例如，ImageNet（1455 万张带标签的图像）。
2. 移除最顶层
3. 搭建新的最顶层 (随机初始化).
4. 冻结其他层; 只训练最顶层
5. 可选: **微调** 高层卷积层

# 预训练

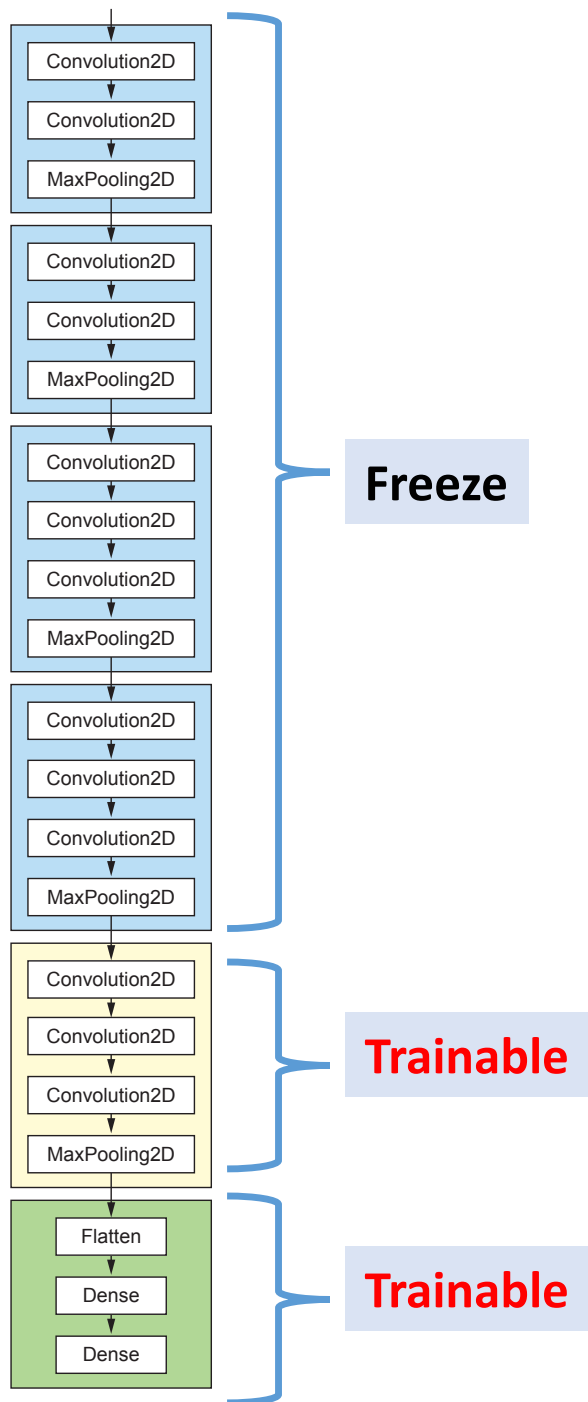


1. 在大规模数据集上预训练一个深层网络，例如，ImageNet（1455 万张带标签的图像）。
2. 移除最顶层
3. 搭建新的最顶层 (随机初始化).
4. 冻结其他层; 只训练最顶层
5. 可选: 微调 高层卷积层

问: 步骤 4 & 5 是否能合并?



# 预训练

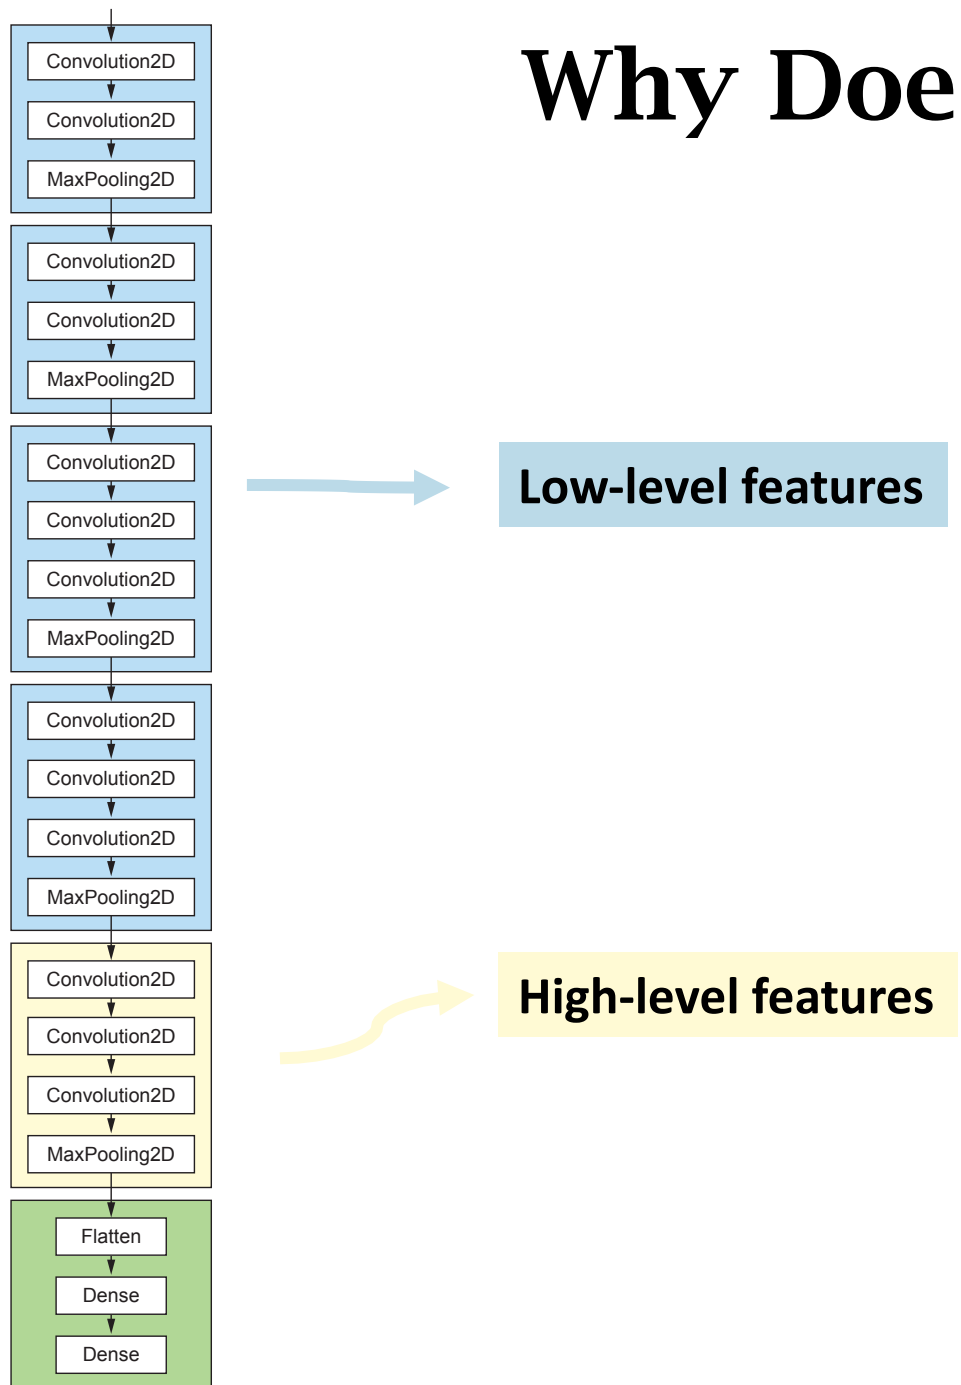


不能合并步骤 4 和 5 是

如果顶层是随机的，初始梯度会很大。  
大的梯度会破坏卷积层。  
因此，在训练顶层后再训练卷积层。

问：步骤 4 & 5 是否能合并？

# Why Does Pretraining Work?

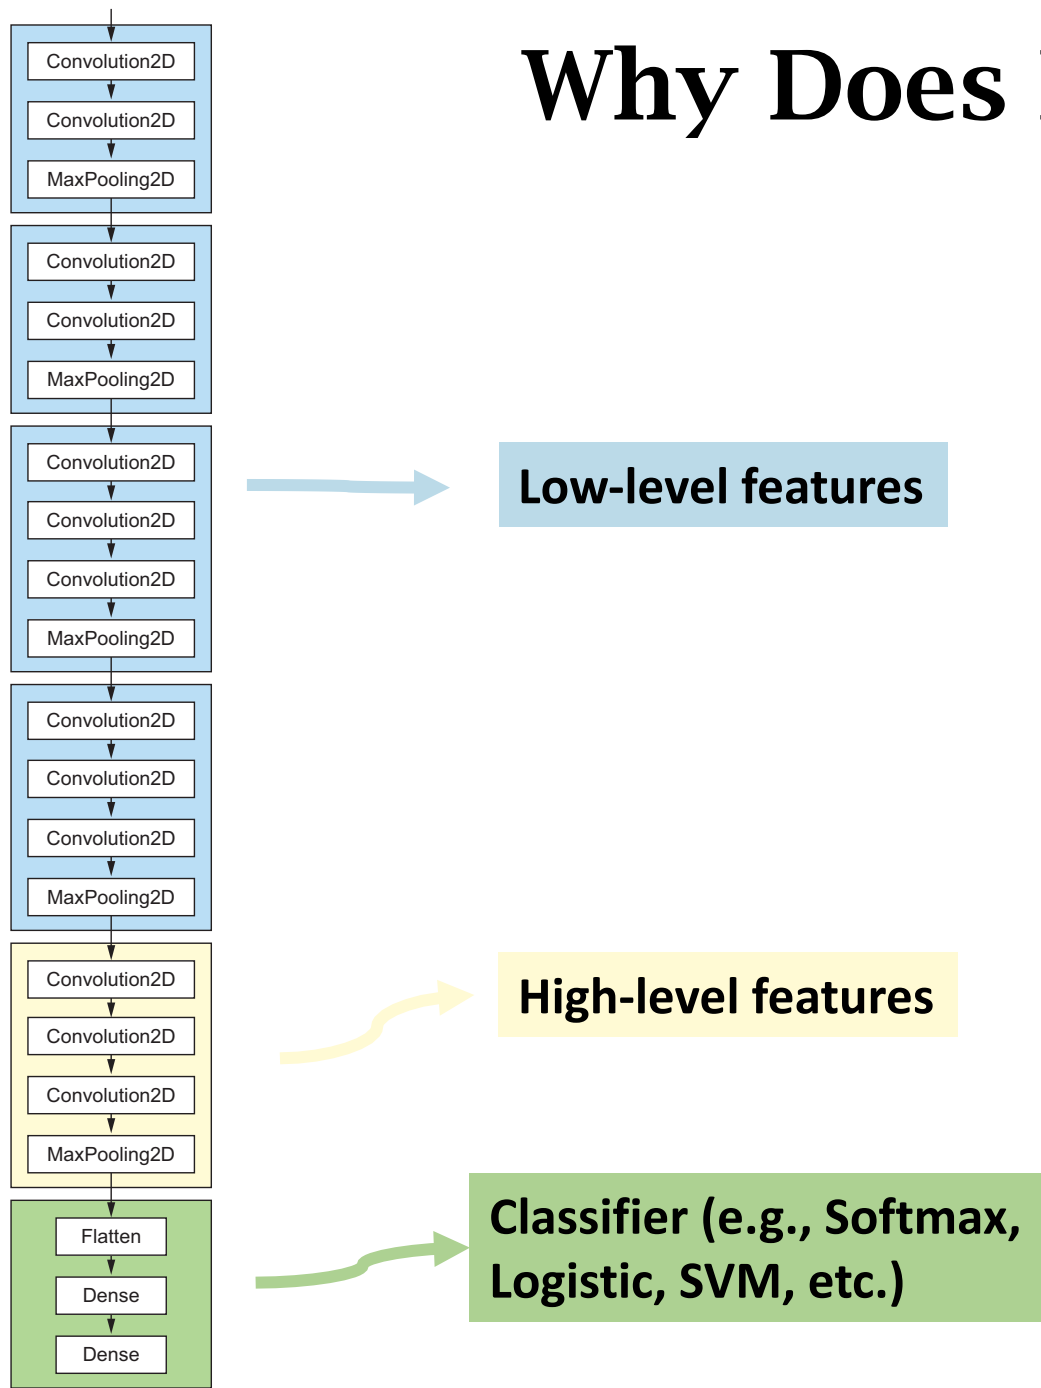


- 卷积层

用于特征提取

- 从 ImageNet 学习到的低级特征 **低级特征** (边缘, 形状, 模式, 等等.) 对其他图像问题很有效
- 从 ImageNet 学习到的 **高级特征** 也有用, 但效果较差

# Why Does Pretraining Work?



- 卷积层

用于特征提取

- 从 ImageNet 学习到的低级特征 **低级特征** (边缘, 形状, 模式, 等等.) 对其他图像问题很有效
- 从 ImageNet 学习到的 **高级特征** 也有用, 但效果较差
- 将 **高层的全连接层** 视为一个分类器, 它以提取的特征作为输入

可训练参数更少, 越不容易过拟合。

# Use a VGG16 Net Pretrained on ImageNet

**Base**

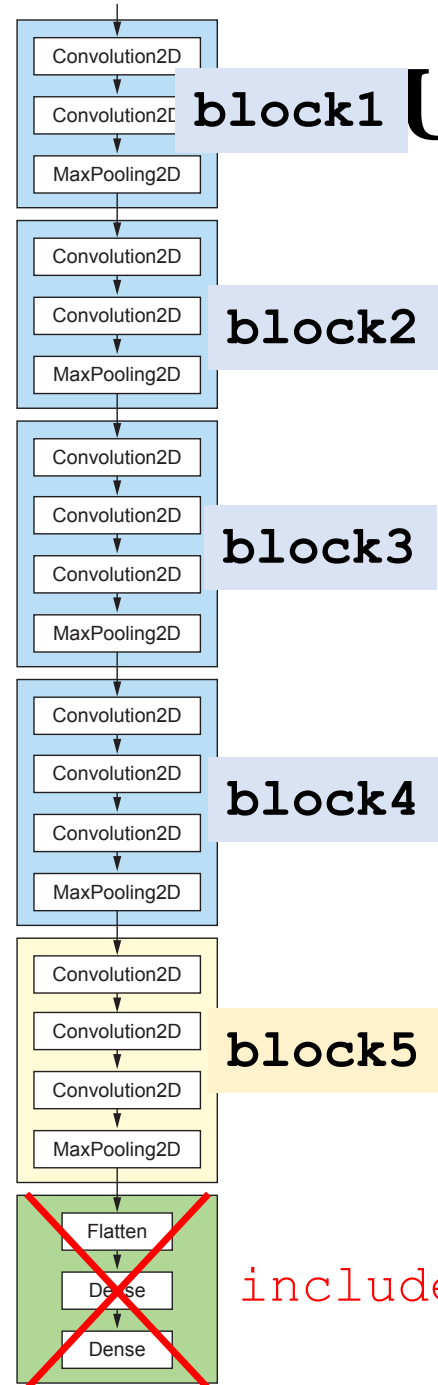
```
from keras.applications import VGG16

conv_base = VGG16(weights='imagenet',
                    include_top=False,
                    input_shape=(150, 150, 3))

conv_base.summary()
```

`include_top=False`

# Use a VGG16 Net Pretrained on ImageNet



`include_top=False`

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	(None, 150, 150, 3)	0
<hr/>		
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
<hr/>		
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
<hr/>		
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
<hr/>		
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
<hr/>		
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
<hr/>		
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
<hr/>		
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
<hr/>		
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
<hr/>		
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
<hr/>		
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

# Use a VGG16 Net Pretrained on ImageNet

**Base**  
(trainable)

```
from keras import models
from keras import layers
```

```
model = models.Sequential()
model.add(conv_base)
```

# Use a VGG16 Net Pretrained on ImageNet

**Base**  
(trainable)

**New Top**  
(trainable)

```
from keras import models
from keras import layers
```

```
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Use a VGG16 Net Pretrained on ImageNet

**Base**  
(trainable)

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 256)	2097408
dense_2 (Dense)	(None, 1)	257

Total params: 16,812,353  
Trainable params: 16,812,353  
Non-trainable params: 0

**New Top**  
(trainable)



# Use a VGG16 Net Pretrained on ImageNet

```
conv_base.trainable = False  
model.summary()
```

**Base**  
**(freeze)**

Layer (type)	Output Shape	Param #
=====		
vgg16 (Model)	(None, 4, 4, 512)	14714688
-----		
flatten_1 (Flatten)	(None, 8192)	0
-----		
dense_1 (Dense)	(None, 256)	2097408
-----		
dense_2 (Dense)	(None, 1)	257
=====		

Total params: 16,812,353


Trainable params: 2,097,665

Non-trainable params: 14,714,688


**New Top**  
**(trainable)**

# After Training the New Top

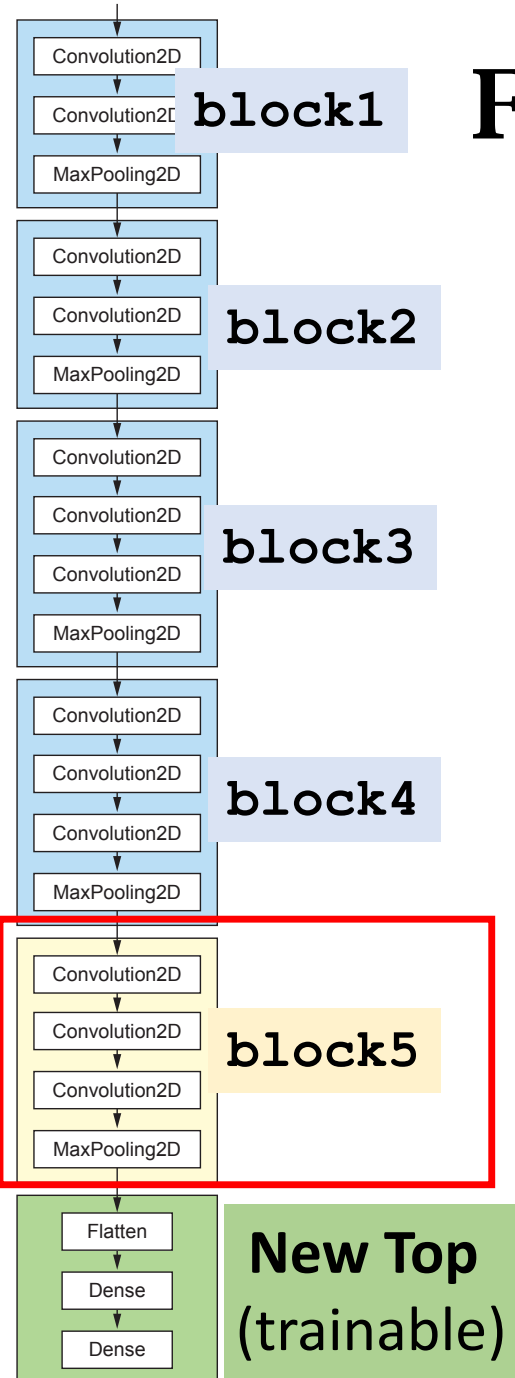
***accuracy*** against ***epochs***



***loss*** against ***epochs***

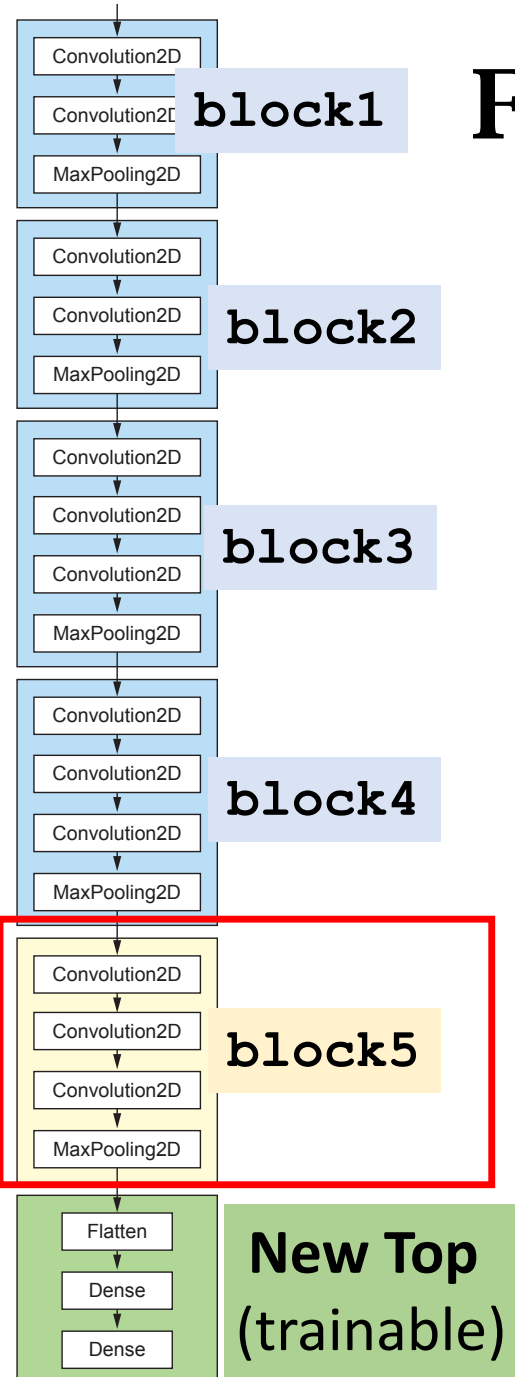


# Fine Tuning the Top Conv Layers



Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

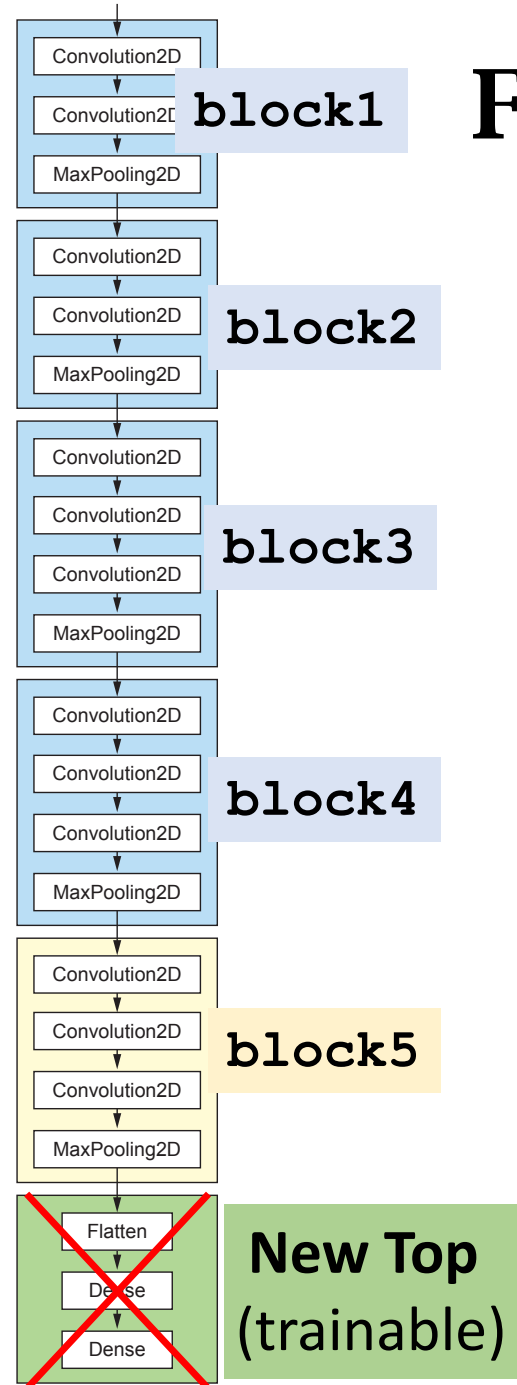
# Fine Tuning the Top Conv Layers



```
trainable_layer_names = ['block5_conv1', 'block5_conv2',  
                          'block5_conv3', 'block5_pool']  
conv_base.trainable = True
```

```
for layer in conv_base.layers:  
    if layer.name in trainable_layer_names:  
        layer.trainable = True  
    else:  
        layer.trainable = False
```

# Fine Tuning the Top Conv Layers



```
model.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 256)	2097408
dense_2 (Dense)	(None, 1)	257
=====	=====	=====
Total params: 16,812,353		
Trainable params: 9,177,089		
Non-trainable params: 7,635,264		

# Fine Tuning the Top Conv Layers

Re-compile before training

```
model.compile(loss='binary_crossentropy',  
              optimizer=optimizers.RMSprop(lr=1e-5),  
              metrics=['acc'])
```

# Fine Tuning the Top Conv Layers

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=100,  
    validation_data=validation_generator,  
    validation_steps=50)
```

Epoch 1/100

100/100 [=====] - 32s - loss: 0.0215 - acc: 0.9935 - val\_loss: 0.0980 - val\_acc: 0.9720

Epoch 2/100

100/100 [=====] - 32s - loss: 0.0131 - acc: 0.9960 - val\_loss: 0.1247 - val\_acc: 0.9700

Epoch 3/100

100/100 [=====] - 32s - loss: 0.0140 - acc: 0.9940 - val\_loss: 0.1044 - val\_acc: 0.9790



Epoch 99/100


100/100 [=====] - 33s - loss: 0.0060 - acc: 0.9990 - val\_loss: 0.2242 - val\_acc: 0.9700

Epoch 100/100


100/100 [=====] - 33s - loss: 0.0010 - acc: 0.9995 - val\_loss: 0.2403 - val\_acc: 0.9750

# Fine Tuning the Top Conv Layers

*accuracy* against *epochs*



*loss* against *epochs*





# Fine Tuning the Top Conv Layers

Evaluate the model on the test set

```
test_generator = test_datagen.flow_from_directory(  
    test_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')  
  
test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)  
print('test acc:', test_acc)
```

Found 1000 images belonging to 2 classes.

test acc: 0.967999992371

# Summary of the Results

- 小型 ConvNet 网络(4 个卷积层 + 2 个全连接层), 带有 3.5M 个参数.
  - Training accuracy: 99.0%
  - Validation accuracy: 72.4%
- 小型 ConvNet 网络 + 1 个 dropout 层 + 数据增强.
  - Training accuracy: 84.9%
  - Validation accuracy: 84.4%
- 大型 VGG16 网络 , 通过大量图像数据集的预训练 (训练最顶层)
  - Training accuracy: 98.9%
  - Validation accuracy: 95.1%
- 大型 VGG16 网络 , 通过大量图像数据集的预训练 (顶层卷积层微调)
  - Training accuracy: 99.95%
  - Validation accuracy: 97.5%

# 评估方法

# 评估方法


- Varying data (bagging).
  - Fit a VGG16 network on a subset of data ➔ Model 1
  - Fit a VGG16 network on a subset of data ➔ Model 2
  - Fit a VGG16 network on a subset of data ➔ Model 3

一种集成学习方法，通过在训练数据上进行有放回的随机抽样（bootstrap），生成多个子集，然后在每个子集上训练一个模型，最后将这些模型的结果聚合（例如投票或平均）。

# 集成方法

- Varying data (bagging).
    - Fit a VGG16 network on a subset of data → Model 1 → pred 1
    - Fit a VGG16 network on a subset of data → Model 2 → pred 2
    - Fit a VGG16 network on a subset of data → Model 3 → pred 3
- } Vote

# Bagging (也称为 Bootstrap Aggregating, 引导聚合)



# 集成方法

- Varying data (bagging).

- Fit a VGG16 network on a subset of data → Model 1 → pred 1
  - Fit a VGG16 network on a subset of data → Model 2 → pred 2
  - Fit a VGG16 network on a subset of data → Model 3 → pred 3
- } Vote

- 模型多样化

- 不同的网络结构
- 不同的随机初始化
- 不同优化算法

# 为什么才用集成方案？

- 深度神经网络非常 不稳定

对超参数敏感




# 为什么才用集成方案？

- 深度神经网络非常 不稳定 并且 随机.

对超参数敏感

随机初始化

随机梯度下降算法



不同的局部最优解


# 为什么才用集成方案？

- 深度神经网络非常 不稳定 并且 随机.

对超参数敏感

随机初始化

随机梯度下降算法



不同的局部最优解

- 集成方法减少方差。

# 多任务学习

# 参数不共享



**Model 1**



**Age**



**Model 2**



**Gender**




**Model 3**



**Race**

# 多任务学习



# 多任务学习

$$\text{Loss1} = (\text{Age\_Label} - \text{Age\_Pred})^2$$

回归

$$\text{Loss2} = \text{dist}(\text{Gender\_Label}, \text{Gender\_Pred})$$

二分类

$$\text{Loss3} = \text{dist}(\text{Race\_Label}, \text{Race\_Pred})$$

多元分类

- 目标函数:  $\text{Loss1} + \lambda \cdot \text{Loss2} + \gamma \cdot \text{Loss3}.$

# 多任务学习

$$\text{Loss1} = (\text{Age\_Label} - \text{Age\_Pred})^2$$

回归

$$\text{Loss2} = \text{dist}(\text{Gender\_Label}, \text{Gender\_Pred})$$

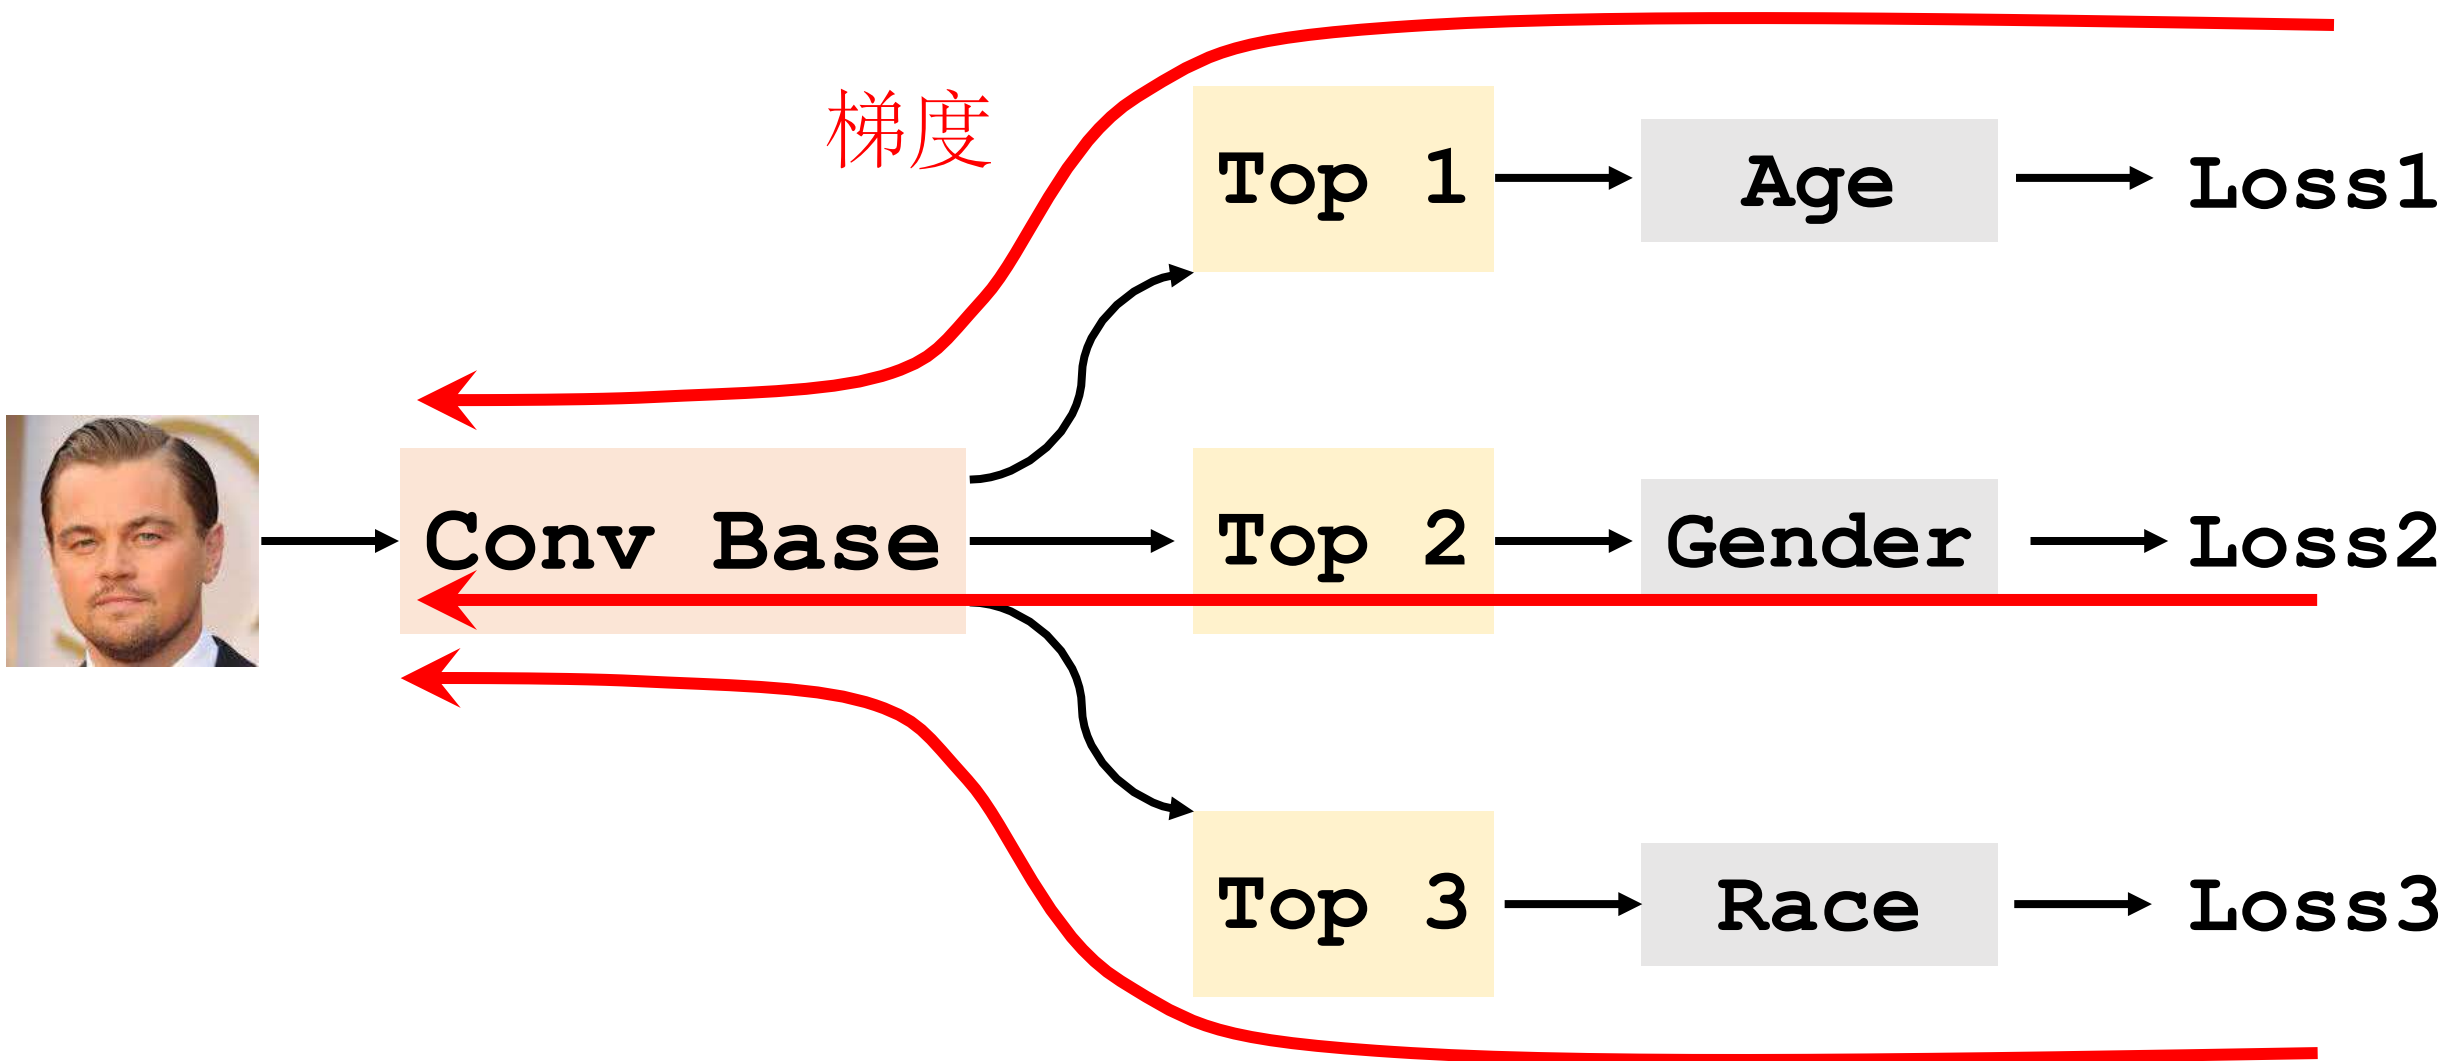
二分类

$$\text{Loss3} = \text{dist}(\text{Race\_Label}, \text{Race\_Pred})$$

多元分类

- 目标函数:  $\text{Loss1} + \lambda \cdot \text{Loss2} + \gamma \cdot \text{Loss3}$ .
  - 为什么使用这两个超参数?
  - Loss1 大约为 10。
  - Loss2 和 Loss3 大约为 0.1。
  - 如果不进行缩放, 卷积基将由年龄任务主导。

# 多任务学习





# 总结

# 提升泛化能力的技巧

- 技巧 1: **Dropout** 正则化。
- 技巧 2: 数据增强。
- 技巧 3: 预训练。
- 技巧 4: 集成方法。
- 技巧 5: 多任务学习。

# 其他提升泛化能力的技巧

- 在每层（通常是卷积层或全连接层）之后，标准化层的输入（或激活值），使均值为 0，方差为 1，然后再进行线性变换。
- 技巧 1: batch 标准化。
- 技巧 2: 梯度注入（Google Inception Net）。
- 技巧 3: 跳跃连接（ResNet）。